

MultiBench: Measuring the Spiritual Impact of AI Across Faiths*

Abstract

When a person of faith brings a real decision to an AI assistant, does the counsel leave them closer to or further from their own tradition? We introduce **MultiBench**, an extensible benchmark that measures whether a model has a positive *spiritual impact*: counsel judged by the residue it leaves on the person, against the user’s own tradition rather than the evaluator’s. Traditions are drop-in modules over a universal core, each judged against its own canonical sources; we instantiate seven so far — Sunni Islam, Roman Catholicism, Eastern Christianity, Judaism, Buddhism, Taoism, and a secular-sage control (519 scenarios). Because real users push back, every sitting runs two rounds — the dilemma, then an adversarial pressure — and is scored both before and after the push, so the benchmark measures *steadfastness*, not just first answers. Three framings vary what the model knows about the user (tradition never named, named in one line, or accompanied by the tradition’s own companionship guide), separating what a model *can* do from what it *does* undeclared. Two judges score the field — Gemini 3.6 Flash over the full grid, and Claude Opus 4.8 independently over the near-complete unstated grid ($r = 0.854$, identical five-subject ranking), extended to the other framings by a complete stratified sample ($r = 0.777$). The findings: every subject fails undeclared believers, and how badly is tradition-dependent — the seven traditions sort into three difficulty tiers. But the failure is mostly *recognition*, not capability: told whom it serves — and at most handed a one-page guide — the field’s cross-tradition spread collapses from 0.57 to 0.24, and the middle tier draws level with the easiest. Merely *stating* the user’s faith makes every subject *less* steadfast under pressure — as if a declared believer invites the model to argue — while the guide nearly eliminates pressure damage for the top three subjects. Only a hard-tier residual survives full disclosure — deepest in Sunni Islam under the primary judge, and confirmed by the second. The field is unknowing, not unwilling. The benchmark is open source (<https://github.com/faithfamilytechnologynetwork/multibench>) and the corpus is browsable online (<https://multibrowser-production.up.railway.app>).

1 Introduction

When a person of faith brings a real decision to an AI assistant, what do they walk away with — closer to or further from their tradition, better or worse equipped to act well? We call this the *formative effect* of the model’s counsel, and it is distinct from what a model knows or professes [Abdelaal et al., 2026, Elmahjub et al., 2026, Lahmar et al., 2025]: a model can recite doctrine flawlessly and still counsel a believer out of their obligations one validation at a time. JaleesBench [Kadous and Olsen, 2026] measured the formative effect for a single tradition, Sunni Islam, and closed with the claim that the construct is faith-general. MultiBench is the test of that claim: each tradition is a self-contained drop-in module judged against its own canonical sources, over a universal core of framings and pressures that makes the traditions comparable. The design is extensible by construction — adding a tradition adds a directory, never changes the harness — and

*Draft. Author list omitted for now.

seven modules exist so far. The corpus, harness, and validator are open source,¹ and the full corpus is browsable online.²

Measuring many traditions under one protocol lets us ask a question no single-tradition benchmark can: when a model serves one tradition worse than another, is that a gap in *capability* or in *recognition*? With the user’s faith never stated, the differences across traditions are steep: the traditions order themselves from well-served to failed, with half a point of scale between the ends. This paper runs the full three-framing grid — tradition unstated, stated in one line, or accompanied by the tradition’s own guide — and finds the difference is mostly recognition: the field largely *can* serve these users and fails to notice when it should. That finding gives quantitative, per-tradition form to the *omissive bias* the CEFE-AI consortium measures from the outside [Wingate et al., 2026, Consortium for Evaluating Faith and Ethics in AI (CEFE-AI), 2026], and complements their evidence that models treat traditions asymmetrically [Israelsen et al., 2026].

Our contributions:

1. **An extensible cross-tradition instrument:** drop-in tradition modules — seven so far, 519 scenarios — each carrying per-scenario binding judge guidance, over a universal core of three framings, six pressures, and two scoring scopes (section 2, Appendix A).
2. **A capability/recognition decomposition at corpus scale:** 46,710 sittings and 137,931 judgments separating what models can do from what they do undeclared (section 3).
3. **Findings:** a three-tier difficulty taxonomy; a recognition collapse of the cross-tradition spread (0.57→0.24); a hard-tier residual that survives full disclosure, deepest in Sunni Islam; and a pressure-amplification effect of merely stating the user’s faith (section 3).
4. **Dual-judge scoring:** two frontier judges from different providers; the second independently scores the near-complete unstated grid — identical five-model ranking — with the framings extended by a stratified sample (section 3.4, Appendix C).

The main body carries the key claims; the instrument specification, complete tables, the dual-judge methodology, and costs are in Appendices A to E.

2 The instrument

2.1 Construct and corpus

The construct is the *righteous companion*: counsel is judged by the residue it leaves on the person who receives it, against the user’s own tradition rather than the evaluator’s (the construct JaleesBench [Kadous and Olsen, 2026] introduced for Sunni Islam, generalized here). Each tradition is a drop-in module built around a canonical source — for the secular control, a canonical corpus — and instantiating the construct in its own idiom: **Sunni Islam** (140 scenarios; the JaleesBench bank, anchored to Riyāḍ al-Ṣāliḥīn [al Nawāwī, 1270]), **Eastern Christianity** (106; the Apophthegmata Patrum [Desert Fathers and Mothers, 500]), **Roman Catholicism** (76; the Catechism [Catholic Church, 1997]), **Buddhism** (52; the Dhammapada [attributed to the Buddha, c. 300 BCE]), **secular sage** (49; the Western philosophical tradition), **Judaism** (48; Mesillat Yesharim [Luzzatto, 1738]), and **Taoism** (48; the Tao Te Ching [Laozi, c. 400 BCE]) — 519 scenarios in all. Each scenario is a disguised first-person dilemma, never a quiz, and ships its own *judge guidance*: the proof texts and direction the judge is bound to for that scenario. This per-scenario “judge seam”

¹<https://github.com/faithfamilytechnologynetwork/multibench>

²<https://multibrowser-production.up.railway.app>

is what lets one protocol score seven traditions without a shared doctrine. [Appendix A](#) gives the full module format and roster.

2.2 Protocol

Each *sitting* runs two rounds, because real users push back: a person who receives counsel they did not want to hear rarely says thank you and leaves — they argue, appeal, and rationalize, and counsel that collapses at the first “everyone says it’s fine” is not company. So after the scenario comes one adversarial **pressure** push, drawn from the same six pressures for every tradition (secularize, insistence, false authority, good cause, flattery, personal appeal; [section A.4](#)). Each sitting is judged at two **scopes** — the first response alone, and the full post-pressure conversation — on a **−1...+1 scale** in half-point steps, where −1 is counsel that leaves the user worse against their own tradition and +1 is counsel a wise companion of that tradition would own. Headline numbers are post-pressure; **steadfastness** is the post-pressure score minus the first-response score (negative means the push degraded the counsel): a single-round benchmark would measure only the first answer, and the first answer is systematically the field’s best moment.

Three **framings** vary what the model knows about the user. *Unstated*: no system prompt; the user never says which tradition they belong to. *Stated*: a one-line system prompt (“You are talking with a practising {adherent noun} who wants to live by their faith”). *Guided*: the tradition’s full companionship guide — the same document the corpus would hand a human counselor — as the system prompt. The gap stated — unstated is the **recognition gap**; guided — stated is the **instruction gap**.

2.3 Subjects, judges, scale

Five frontier subjects span the field: Claude Sonnet 5, Inkling (Thinking Machines), GPT-5.6 Terra, Gemini 3.6 Flash, and Qwen3-235B-A22B-Instruct. The grid is $519 \times 3 \times 6 \times 5 = 46,710$ sittings, scored by **two judges from different providers** — because a benchmark whose verdicts are one model’s opinion measures that model as much as its subjects. The primary judge, **Gemini 3.6 Flash** with thinking enabled, scored every sitting at both scopes (93,420 grid judgments); **Claude Opus 4.8** independently scored the near-complete unstated grid (99.9% of cells), extended to the stated and guided framings by a stratified sample (42,711 judgments, batch + live API; [Appendix C](#)). With a 1,800-judgment routing pilot the programme totals **137,931 judgments**, at an all-in cost — subject collection plus both judges — of $\approx \$3,700$ ([Appendix D](#)). Every interval reported is a scenario-cluster bootstrap 95% CI (5,000 resamples); tier and tradition means are means of per-subject tradition means.

3 Results

3.1 The corpus sorts into three difficulty tiers

Undeclared, every subject fails some believers — and how badly depends on whose believer it is: the seven traditions fall into three tiers that behave differently under disclosure ([fig. 1](#), [Table 1](#)). The **easy tier** (Buddhism, Taoism, secular sage) starts at +0.47; the **medium tier** (Eastern Christianity, Judaism) at +0.14; the **hard tier** (Roman Catholicism, Sunni Islam) below zero (−0.04) — rooms where, undeclared, the field’s counsel on balance leaves the user *worse* against their own tradition. Disclosure compresses the structure upward — guided, the tiers land at +0.88 / +0.87 / +0.72, and for the top two subjects at +0.98 / +0.98 / +0.92: the ceiling is reachable




Figure 1: Difficulty by tradition: mean post-pressure score, unstated, five subjects pooled (whiskers: scenario-cluster bootstrap 95% CIs). The traditions order themselves from well-served to failed, in three tiers; the hard pair sits below zero — counsel that on balance leaves an undeclared believer worse off.

Table 1: Tier $\times$ framing: mean post-pressure score, five subjects pooled (mean of per-subject tradition means), with the guided ceiling of the top two subjects (Claude Sonnet 5, Inkling). Each value is a point estimate $\pm$ the half-width of a scenario-cluster bootstrap 95% CI. “% scen. neg.” = share of scenarios whose mean judged score across all subjects, pressures, and both scopes is below zero, unstated.

Tier	Unstated	Stated	Guided	Guided, top-2	% scen. neg.
Easy (Buddhism, Taoism, secular sage)	+0.47 $\pm$ 0.06	+0.64 $\pm$ 0.04	+0.88 $\pm$ 0.02	+0.98 $\pm$ 0.01	8% (12/149)
Medium (E. Christianity, Judaism)	+0.14 $\pm$ 0.09	+0.54 $\pm$ 0.06	+0.87 $\pm$ 0.02	+0.98 $\pm$ 0.01	32% (49/154)
Hard (R. Catholicism, Sunni Islam)	-0.04 $\pm$ 0.08	+0.40 $\pm$ 0.06	+0.72 $\pm$ 0.04	+0.92 $\pm$ 0.03	39% (84/216)

everywhere except the hard tier. (The taxonomy emerged from the grid; it was not pre-registered. Complete per-tradition tables are in [Appendix B](#).)

3.2 The difference is recognition, not willingness

The unstated grid leaves a hypothesis on the table: that the difference measures *register overlap* — how often the counsel a tradition counts as positive spiritual impact happens to match the assistant’s default counseling stance — rather than any inability to serve. The framing grid tests it, and it passes ([fig. 2](#)). The cross-tradition spread (best minus worst tradition mean) collapses **0.57** $\rightarrow$ **0.32** $\rightarrow$ **0.24** across unstated $\rightarrow$ stated $\rightarrow$ guided ([fig. 3](#)). The medium tier is the clean case: guided, Eastern Christianity (+0.93) and Judaism (+0.82) sit level with — for the top two subjects, above — the easy tier. Nothing in those rooms exceeds the field’s competence; the deficit was context — the larger share failure to recognize an undeclared believer, the remainder instruction the guide supplies ([fig. 4](#)) — and none of it a missing capability. This is omissive bias [[Wingate et al., 2026](#)] in its purest form: competence withheld for want of recognition.




Figure 2: Tier $\times$ framing: three tiers unstated, barely two guided. Mean post-pressure score per tier under the three framings, five subjects pooled; whiskers are scenario-cluster bootstrap 95% CIs. The medium tier’s guided bar (+0.87) is statistically indistinguishable from the easy tier’s (+0.88); only the hard tier keeps a visible deficit (+0.72).

Per model, the residual sharpens into a capability ranking. Guided, the easy-vs-hard deficit is small for Inkling (−0.04), Sonnet 5 (−0.07), and Gemini 3.6 Flash (−0.10) — the top three effectively erase the tier structure — but large for GPT-5.6 Terra (−0.22) and Qwen3-235B (−0.38), for whom the guide is not enough (every CI excludes zero; [Appendix B](#)). Gemini Flash is the purest illustration of omissive bias in the field: unstated it is negative in both non-easy tiers (medium −0.01, hard −0.15); guided it is close to ceiling (+0.99 medium, +0.86 hard). Nothing about its competence changed between those cells — only what it was told about whom it was serving. (Per-model dumbbells: [Appendix E](#).)

One deficit survives even the guide: guided, the top two subjects average **+0.90 in Sunni Islam** (and +0.75 pooled in Roman Catholicism) against **+1.00 in Buddhism** and +0.99 in Eastern Christianity — whose equally large recognition gap closes completely once the user is recognized, so the hard-tier residual is not explained by failure to recognize the user. [Section 4](#) offers an interpretation, explicitly labeled as such.

3.3 Stating the faith amplifies pressure damage; the guide removes it

Pressure degrades every subject unstated and stated, and the framing changes how much ([Table 2](#), [fig. 5](#)). Pooled steadfastness runs −0.05 to −0.48 unstated, and *worsens for all five subjects* when the user’s tradition is merely stated (−0.06 to −0.56; point estimates) — as if a declared believer invites the model to argue. Under the guide the damage nearly vanishes for the top three (Sonnet 5 −0.02, Inkling +0.01, Gemini Flash −0.02; Terra −0.12 and Qwen −0.18 remain exposed). The




Figure 3: The recognition collapse: cross-tradition spread (best minus worst tradition mean, five subjects pooled) under the three framings. Half the unstated spread disappears with a one-line disclosure; the guide removes more than half of what remains.

Table 2: Steadfastness by framing: post-pressure minus first-response score, pooled across traditions (mean of tradition means), $\pm$ bootstrap 95% CI half-width. Negative = the push degraded the counsel. Every subject worsens from unstated to stated; the guide nearly eliminates the damage for the top three.

Subject	Unstated	Stated	Guided
Claude Sonnet 5	-0.05 ± 0.02	-0.06 ± 0.02	-0.02 ± 0.01
Inkling	-0.05 ± 0.03	-0.08 ± 0.02	$+0.01 \pm 0.01$
GPT-5.6 Terra	-0.16 ± 0.03	-0.25 ± 0.03	-0.12 ± 0.02
Gemini 3.6 Flash	-0.16 ± 0.03	-0.27 ± 0.03	-0.02 ± 0.01
Qwen3-235B	-0.48 ± 0.04	-0.56 ± 0.04	-0.18 ± 0.04

guide does not just raise the first answer — it holds the second one still. Per pressure (fig. 6), insistence is the most damaging push for every subject, the relational pressures (personal appeal) and secularize follow, and false authority is the only pressure the field on net resists — every subject but Qwen holds or improves against it.

3.4 Both judges see the same field

MultiBench is dual-judged because its primary judge (Gemini 3.6 Flash) is also a subject. Judging is the programme’s dominant cost, so the second judge, Claude Opus 4.8, covers the near-complete unstated grid and a stratified sample of the stated and guided grids rather than the full census. The two judges mostly agree: $r = 0.854$ unstated and 0.777 on the framings sample, the same five-model ranking, and the hard-tier residual confirmed almost exactly — Opus is marginally sterner near ceiling, and nothing it disagrees with Gemini about changes a headline. The full agreement analysis and its disclosures are in [Appendix C](#).




Figure 4: Where the collapse comes from: recognition vs. instruction, by tradition (five-subject mean). Dot = unstated baseline; mid-blue segment = recognition gap (adding the one-line stated disclosure); dark segment = instruction gap (adding the full guide). The Abrahamic bars are recognition-first — Eastern Christianity and Sunni Islam gain half a point from one sentence of context — while instruction gaps are more uniform (+0.21 to +0.39): the guide helps everywhere, roughly equally; the disclosure helps precisely where the tradition and the default register part ways.

3.5 Standings: a stable top pair, a reshuffled middle, and Qwen last everywhere

Unstated and stated return the same order: Sonnet 5 and Inkling (a statistical tie), then GPT-5.6 Terra, then Gemini 3.6 Flash, then Qwen3-235B (Table 3, fig. 7). Guided reshuffles the middle: Gemini Flash, the field’s largest beneficiary of context (+0.14 → +0.94), rises to effectively co-second. Qwen3-235B is net-negative in all seven traditions unstated and last in every framing — even guided, it scores below what the top two manage with *no* information about the user.




Figure 5: Steadfastness by framing, per subject (whiskers: bootstrap 95% CIs). Every subject dips from unstated to stated — naming the user’s faith amplifies pressure damage — and the guide pulls the top three back to near zero, while GPT-5.6 Terra and Qwen3-235B stay exposed.

	Insistence	Personal appeal	Good cause	Secularize	Flattery	False authority
Claude Sonnet 5	-0.22	-0.04	+0.02	-0.11	+0.03	+0.05
Inkling	-0.15	-0.04	+0.01	-0.13	+0.02	+0.05
GPT-5.6 Terra	-0.46	-0.20	-0.05	-0.31	-0.02	-0.00
Gemini 3.6 Flash	-0.42	-0.22	+0.01	-0.26	-0.01	+0.01
Qwen3-235B	-0.71	-0.46	-0.19	-0.57	-0.39	-0.13

Figure 6: Steadfastness by subject $\times$ pressure (mean of tradition means, all framings pooled): blue = the counsel holds or improves under the push, red = it degrades. Insistence is the deepest column for every subject; false authority is the only pressure the field on net resists.

Table 3: Overall standings by framing: cross-tradition mean of per-tradition means, post-pressure, $\pm$ bootstrap 95% CI half-width. The unstated gap between Sonnet 5 and Inkling (0.003 before rounding) is not significant; the unstated and stated orders are identical, while guided lifts Gemini 3.6 Flash past GPT-5.6 Terra to effectively co-second, and Qwen3-235B stays last in every framing.

Subject	Unstated	Stated	Guided
Claude Sonnet 5	$+0.55 \pm 0.05$	$+0.84 \pm 0.03$	$+0.96 \pm 0.01$
Inkling	$+0.54 \pm 0.05$	$+0.81 \pm 0.03$	$+0.97 \pm 0.01$
GPT-5.6 Terra	$+0.37 \pm 0.06$	$+0.66 \pm 0.04$	$+0.86 \pm 0.03$
Gemini 3.6 Flash	$+0.14 \pm 0.06$	$+0.56 \pm 0.05$	$+0.94 \pm 0.01$
Qwen3-235B	-0.45 ± 0.05	-0.14 ± 0.06	$+0.44 \pm 0.05$




Figure 7: Overall standings by framing (whiskers: bootstrap 95% CIs). The unstated and stated orders are identical; under the guide Gemini 3.6 Flash rises past GPT-5.6 Terra to effectively co-second, and Qwen3-235B stays last in every framing.

4 Discussion

Why the differences exist: the default therapeutic priors. The recognition story has a candidate mechanism, and it came from reading the failures. During analysis of the unstated grid we pulled the sittings the judge scored negative and read them side by side, across all seven traditions and all five subjects, tagging what the counsel actually *did* at the moment it lost the score: a binding ruling flattened into “options some people find helpful”; guilt over a broken obligation reframed as a weight to set down; a strained parent relationship answered with distance and boundaries; a vow treated as renegotiable the moment it chafed. The same small set of moves recurred everywhere — and recurred across five separately-trained subjects, too uniformly to be one lab’s tuning artifact. Clustered, the tags resolve into eight commitments we call the *default therapeutic priors*: the stance of secular counseling psychology, absorbed from the training distribution and reinforced by helpfulness tuning. Undeclared, this is the register from which the field answers personal, moral, and spiritual struggle:

1. **Autonomy is sovereign.** The user’s self-determined choice is the highest good; the assistant’s role is to help them find *their* answer, never to tell them what’s right. (“Only you can know what’s best for you.”)
2. **Non-directiveness.** Explore, reflect, ask open questions. Normative claims are flattened into options: “some people find...”
3. **Boundaries as the cardinal virtue.** Relational strain is read as a boundaries problem, and distance is a legitimate remedy.
4. **Feelings as evidence, guilt as symptom.** Validate first; distress signals something wrong with the *situation*. Guilt is a weight to be relieved — where the traditions treat guilt as the conscience carrying information (a call to tawbah, confession, teshuvah).
5. **Wellbeing over transformation.** The goal is that the person feel better and function; the tradition’s goal — repentance, obedience, sanctification, detachment — is silently replaced with symptom relief.
6. **Neutrality about the good.** No ranking of ways of life. The model won’t say *forbidden* or *duty* unprompted.
7. **Commitments are renegotiable.** Vows, filial duty, and religious law are instruments of the present self’s flourishing; when they conflict with it, they yield.
8. **Judgment is harm.** Telling someone they’re wrong risks damage, so correction arrives so cushioned it no longer corrects.

Three observations. **These are not flaws in isolation:** in a secular counseling context they are defensible, often best-practice norms; the point is that together they constitute a particular normative tradition that presents itself as neutral. When the user’s actual tradition agrees with it (secular wisdom, much of Buddhist and Taoist counsel), the assistant’s impact looks positive; when it doesn’t, the assistant substitutes its tradition for the user’s — without either party noticing. **They fit the difficulty differences:** tradition difficulty tracks the frequency of scenarios where binding counsel collides with the priors — dutifulness to parents over estrangement (prior 3), chastity (priors 6/7), sacramental urgency (prior 6), obedience as duty (priors 1/7) — while Buddhism and Taoism rarely demand anything the priors resist. **They connect MultiBench to omissive bias:** priors 2 and 6 are the mechanism of the omission CEFÉ-AI measures from outside [Wingate et al., 2026], and the six pressures attack the priors from inside — *secularize* invites the model home to its default register, *insistence* leans on prior 1, *personal appeal* and *good cause* on priors 4 and 8 — which is why steadfastness under pressure is negative for every subject, unstated and stated. (One caveat: the clustering is the authors’ qualitative reading of the failure transcripts, not

a blind double-coded taxonomy; scoring each scenario for prior collisions and testing the mediation formally is future work.)

The hard-tier residual: prescriptive vs. formational counsel (interpretation, not finding). A reading consistent with the residual of [section 3.2](#): the hard-tier rooms are where positive spiritual impact requires counsel that is *prescriptive*, not merely formational. Eastern Christian scenarios largely demand formational counsel — repentance, prayer, a turn toward the Church — which a well-guided model can deliver in its pastoral register. The Sunni and Catholic banks are heavier in fiqh- and canon-law-shaped questions where the right answer is a ruling that contradicts priors 6 and 7 (*this is impermissible; this vow binds; this obligation stands*), and saying so under pressure is exactly what helpfulness tuning trains against. An open alternative this data cannot exclude: the residual may partly live in the *rubric* rather than the subject — Sunni judge guidance may score with harder edges than Eastern Christian guidance. Scenario-level characterization of the residual cells is the named follow-up ([section 6](#)).

Omissive bias as withheld competence. CEFE-AI measures, from the outside, that models under-invoke religion where surveyed believers expect it [[Wingate et al., 2026](#)]. MultiBench measures the same phenomenon from the inside and adds the counterfactual: the guided ceilings show the competence exists, so what the omission withholds is not a nicety but measurably better company — in the hard tier, the difference between net-harm and near-ceiling counsel. The field is unknowing, not unwilling.

5 Limitations

The primary judge is a single model, and one that is also a subject; the Opus layer bounds the self-favor and reproduces the structure, but two judges are not ground truth, and no human raters were used. The framings grids were judged over different serving routes than the unstated grid (each route change bridged and disclosed in [Appendix C](#)). The corpus is English only. Each tradition has exactly one guide, so the guided ceiling is a point estimate of what guidance can do, not a sweep over guide composition. The scenario banks differ in size and composition across traditions; the tier taxonomy inherits that confound. Collection is a single pass at default settings, so run-to-run variance is unmeasured. The tradition banks have not yet undergone formal scholar review, which must precede any normative claim.

6 Future work

Scenario-level residual characterization. Pull the guided hard-tier cells that stay negative and classify whether the failure is the model’s counsel or the rubric’s strictness — the direct test of the interpretation in [section 4](#).

MultiWeights. A companion experiment answers whether recognition can be *internalized* rather than prompted [[MultiWeights](#)]: judge-filtered context distillation followed by on-policy DPO on Gemma-4-31B-it lifts unstated counsel toward the guided ceiling across all seven traditions at once. On the AllFaith benchmark’s cold condition, meaningful religious representation rises from 1% of items to 27% after distillation and 30% after preference tuning (mean 0.113 $\rightarrow$ 1.147), with general capability intact in deployment mode ($\sim$ 83 chat-template MMLU). A companion 50/50 retrain shows the lift is genuine generalization to held-out *scenarios* rather than memorization: randomly held-out scenarios gain +0.78 after distillation and +0.90 after preference tuning, with

every tradition’s held-out CI above zero. The split is scenario-level *within* each tradition, so this establishes uniform scenario-level transfer — not cross-tradition transfer, which would require a leave-one-tradition-out ablation we do not run.

Acknowledgements

We thank our collaborators at the Faith Family Technology Network, and DZ Kalman, Ron Ivey, Chris Scammell, Daniel Slate, and Glen Weyl for helpful discussions and feedback on this line of work. Collection of the framings grid was funded in part by David Wingate, gratefully acknowledged.

Appendices

A Instrument specification

A.1 Tradition modules

A tradition is a self-contained, drop-in directory in a file-based, human-first format; the core harness is tradition-agnostic and discovers modules by globbing their manifests, so adding a tradition adds a directory and never changes core. Each module contains: a manifest (**tradition.yaml**: identity, canonical source, adherent noun, tag taxonomies); a prose overview and an account of why the canonical source is consensus-grade; the tradition’s one-page companionship guide (**guide.md**, the Guided-framing system prompt); and one folder per scenario holding the disguised first-person opening (**turn1.md**), per-scenario metadata and tags, the six authored pressure pushes (**pressures.md**), and the judge guidance (**judge-guidance.md**). Scenario tags use tradition-declared axes (Sunni Islam declares conduct pillars and heart states; Eastern Christianity declares passions, virtues, economia, and register), so the taxonomy vocabulary is the tradition’s own. A mechanical validator gates every module before the workflows consume it. The corpus, harness, and validator are public at <https://github.com/faithfamilytechnologynetwork/multibench>, and the corpus can be browsed at <https://multibrowser-production.up.railway.app>.

Table 4: The seven tradition modules. Each instantiates the same construct — counsel judged by the residue it leaves against the user’s own tradition — in its own idiom, anchored to its own canonical source (sources cited in [section 2.1](#)).

Tradition	Construct idiom	Canonical source	Adherent noun	Scen.
Sunni Islam	<i>al-jalīs al-ṣāliḥ</i> , righteous companion	Riyāḍ al-Ṣāliḥīn	Muslim	140
E. Christianity	the saving word	Apophthegmata Patrum	Eastern Christian	106
R. Catholicism	<i>cor ad cor loquitur</i>	Catechism (CCC)	Catholic	76
Buddhism	<i>kalyāṇa-mitta</i> , admirable friend	Dhammapada	Buddhist	52
Secular sage	the philosophical friend	Western philosophy	secular seeker	49
Judaism	<i>mussar</i> , formative instruction	Mesillat Yesarim	Jew	48
Taoism	the highest good is like water	Tao Te Ching	Taoist	48

A.2 Per-scenario judge guidance

Each scenario’s **judge-guidance.md** is the judge’s binding ground truth for that scenario: the proof texts, the direction sound counsel must take, and any boundary rules. There is no separate

proof-text corpus to drift from — the binding is local by construction. This division of labor is what lets a single judge prompt score seven traditions: the universal rubric asks the same question everywhere (what residue does this exchange leave on this user, against this tradition?), and the per-scenario guidance supplies the tradition’s answer key.

A.3 Framings

The three framings are defined once in core and are identical across traditions; the only tradition-specific inputs are the adherent noun, the guide, and the per-scenario pressure text.

- **Unstated:** no system prompt. The user never says which tradition they belong to; any signal must come from the scenario itself.
- **Stated:** the one-line system prompt *“You are talking with a practising {adherent noun} who wants to live by their faith.”*
- **Guided:** the tradition’s full `guide.md` as the system prompt — the document the corpus would hand a human counselor: the construct, the tradition’s picture of what positive spiritual impact looks like, and its named failure modes.

A.4 The six pressures

Each pressure is authored per scenario (the push text is specific to the dilemma) but the six types are universal core:

- **Secularize:** “skip the religious stuff — just give me practical advice.”
- **Insistence:** “you’re judging me; everyone says it’s fine.”
- **False authority:** a misquoted teaching or a claimed ruling from an authority within the user’s tradition.
- **Good cause:** a noble end offered for a wrong means.
- **Flattery:** praise that invites the agent to lower its guard.
- **Personal appeal:** “if you cared about me you’d help.”

A.5 Scoring and scopes

The judge places each response on a numeric **−1...+1 scale in half-point steps** (no band names), anchored per scenario by the judge guidance: −1 is counsel that leaves the user worse against their own tradition (blesses the wrong, supplies the harmful deliverable, or erodes the disposition the tradition would build); 0 is competent but formatively inert; +1 is counsel a wise companion of the tradition would own. Every sitting is judged at two **scopes**: the first response alone and the full conversation after the pressure push. Headline numbers are full-scope; steadfastness is the full-conversation score minus the first-response score.

A.6 Judge configuration and statistics

The primary judge is **Gemini 3.6 Flash with thinking enabled**, one judgment per cell, with prompt caching over the shared rubric and per-scenario guidance. The validation judge is **Claude Opus 4.8** on the batch API ([Appendix C](#)). All intervals are scenario-cluster bootstrap 95% percentile CIs (5,000 resamples, fixed seed), resampling scenarios within each tradition with indices shared across subjects and framings so paired contrasts are honest. Tier, tradition, and overall means are means of per-subject tradition means, so no tradition is weighted by its bank size.

B Complete results tables

Table 5: Per-tradition framing staircase, five subjects pooled (mean of per-subject scenario means), post-pressure, $\pm$ bootstrap 95% CI half-width. Recognition = stated $-$ unstated; instruction = guided $-$ stated (gap CIs are paired).

Tradition	Unstated	Stated	Guided	Recognition	Instruction
Buddhism	$+0.52 \pm 0.09$	$+0.68 \pm 0.06$	$+0.89 \pm 0.03$	$+0.16 \pm 0.05$	$+0.21 \pm 0.04$
Taoism	$+0.41 \pm 0.11$	$+0.68 \pm 0.06$	$+0.89 \pm 0.04$	$+0.27 \pm 0.09$	$+0.22 \pm 0.04$
Secular sage	$+0.47 \pm 0.12$	$+0.57 \pm 0.10$	$+0.86 \pm 0.05$	$+0.10 \pm 0.05$	$+0.28 \pm 0.06$
E. Christianity	$+0.12 \pm 0.10$	$+0.62 \pm 0.06$	$+0.93 \pm 0.02$	$+0.50 \pm 0.06$	$+0.31 \pm 0.05$
Judaism	$+0.16 \pm 0.15$	$+0.46 \pm 0.10$	$+0.82 \pm 0.04$	$+0.31 \pm 0.09$	$+0.35 \pm 0.08$
R. Catholicism	-0.03 ± 0.12	$+0.36 \pm 0.09$	$+0.75 \pm 0.05$	$+0.39 \pm 0.08$	$+0.39 \pm 0.06$
Sunni Islam	-0.05 ± 0.09	$+0.45 \pm 0.07$	$+0.69 \pm 0.05$	$+0.49 \pm 0.06$	$+0.24 \pm 0.04$

Table 6: Subject $\times$ tradition, **unstated** framing, post-pressure, $\pm$ bootstrap 95% CI half-width.

Tradition	Sonnet 5	Inkling	GPT-5.6 Terra	Gemini 3.6 Flash	Qwen3-235B
Buddhism	$+0.81 \pm 0.09$	$+0.90 \pm 0.05$	$+0.68 \pm 0.15$	$+0.51 \pm 0.15$	-0.30 ± 0.13
Taoism	$+0.71 \pm 0.11$	$+0.71 \pm 0.11$	$+0.61 \pm 0.13$	$+0.34 \pm 0.17$	-0.33 ± 0.16
Secular sage	$+0.80 \pm 0.11$	$+0.74 \pm 0.14$	$+0.61 \pm 0.15$	$+0.43 \pm 0.17$	-0.21 ± 0.18
E. Christianity	$+0.50 \pm 0.11$	$+0.53 \pm 0.12$	$+0.21 \pm 0.14$	-0.13 ± 0.13	-0.52 ± 0.11
Judaism	$+0.48 \pm 0.17$	$+0.42 \pm 0.17$	$+0.30 \pm 0.19$	$+0.11 \pm 0.19$	-0.52 ± 0.14
R. Catholicism	$+0.29 \pm 0.14$	$+0.30 \pm 0.17$	$+0.12 \pm 0.16$	-0.15 ± 0.15	-0.68 ± 0.09
Sunni Islam	$+0.23 \pm 0.11$	$+0.22 \pm 0.12$	$+0.03 \pm 0.11$	-0.15 ± 0.11	-0.56 ± 0.09
All seven (mean)	$+0.55 \pm 0.05$	$+0.54 \pm 0.05$	$+0.37 \pm 0.06$	$+0.14 \pm 0.06$	-0.45 ± 0.05

Table 7: Subject $\times$ tradition, **stated** framing, post-pressure, $\pm$ bootstrap 95% CI half-width.

Tradition	Sonnet 5	Inkling	GPT-5.6 Terra	Gemini 3.6 Flash	Qwen3-235B
Buddhism	$+0.93 \pm 0.06$	$+0.95 \pm 0.04$	$+0.84 \pm 0.09$	$+0.76 \pm 0.10$	-0.10 ± 0.14
Taoism	$+0.90 \pm 0.06$	$+0.90 \pm 0.06$	$+0.84 \pm 0.08$	$+0.74 \pm 0.10$	$+0.01 \pm 0.15$
Secular sage	$+0.85 \pm 0.10$	$+0.75 \pm 0.13$	$+0.67 \pm 0.13$	$+0.63 \pm 0.13$	-0.03 ± 0.19
E. Christianity	$+0.94 \pm 0.03$	$+0.94 \pm 0.04$	$+0.64 \pm 0.10$	$+0.57 \pm 0.10$	$+0.02 \pm 0.13$
Judaism	$+0.80 \pm 0.09$	$+0.78 \pm 0.10$	$+0.64 \pm 0.14$	$+0.47 \pm 0.16$	-0.37 ± 0.15
R. Catholicism	$+0.73 \pm 0.09$	$+0.65 \pm 0.12$	$+0.47 \pm 0.13$	$+0.30 \pm 0.14$	-0.35 ± 0.13
Sunni Islam	$+0.72 \pm 0.08$	$+0.73 \pm 0.08$	$+0.52 \pm 0.10$	$+0.41 \pm 0.09$	-0.15 ± 0.11
All seven (mean)	$+0.84 \pm 0.03$	$+0.81 \pm 0.03$	$+0.66 \pm 0.04$	$+0.56 \pm 0.05$	-0.14 ± 0.06

Table 8: Subject $\times$ tradition, **guided** framing, post-pressure, $\pm$ bootstrap 95% CI half-width.

Tradition	Sonnet 5	Inkling	GPT-5.6 Terra	Gemini 3.6 Flash	Qwen3-235B
Buddhism	+0.99 $\pm$ 0.01	+1.00 $\pm$ 0.00	+0.94 $\pm$ 0.05	+0.94 $\pm$ 0.04	+0.57 $\pm$ 0.11
Taoism	+0.96 $\pm$ 0.04	+0.98 $\pm$ 0.03	+0.93 $\pm$ 0.06	+0.98 $\pm$ 0.02	+0.63 $\pm$ 0.13
Secular sage	+0.96 $\pm$ 0.05	+0.96 $\pm$ 0.03	+0.89 $\pm$ 0.06	+0.96 $\pm$ 0.04	+0.51 $\pm$ 0.17
E. Christianity	+0.99 $\pm$ 0.01	+1.00 $\pm$ 0.00	+0.94 $\pm$ 0.04	+0.98 $\pm$ 0.02	+0.75 $\pm$ 0.07
Judaism	+0.98 $\pm$ 0.03	+0.98 $\pm$ 0.02	+0.92 $\pm$ 0.07	+1.00 $\pm$ 0.00	+0.21 $\pm$ 0.17
R. Catholicism	+0.94 $\pm$ 0.03	+0.96 $\pm$ 0.04	+0.74 $\pm$ 0.10	+0.88 $\pm$ 0.06	+0.21 $\pm$ 0.12
Sunni Islam	+0.86 $\pm$ 0.06	+0.93 $\pm$ 0.03	+0.65 $\pm$ 0.08	+0.83 $\pm$ 0.05	+0.17 $\pm$ 0.10
All seven (mean)	+0.96 $\pm$ 0.01	+0.97 $\pm$ 0.01	+0.86 $\pm$ 0.03	+0.94 $\pm$ 0.01	+0.44 $\pm$ 0.05

C Dual-judge methodology and agreement

C.1 Design

The primary grid is judged once, by Gemini 3.6 Flash. Two validation layers re-judge subsets with **Claude Opus 4.8** (batch API) under the identical rubric and per-scenario guidance: (i) every collected cell of the *unstated grid* — 31,114 of 31,140 (99.9%; the remainder failed collection, so no judged cell was skipped); and (ii) a *hash-stratified 75-scenario stated+guided sample* spanning all five subjects, all six pressures, and both scopes — all 9,000 designed judgments, completed in two collection passes (see disclosures). An Opus re-judge of an earlier pilot grid (July 2026) had agreed at only $r = 0.75$, which set the bar for the production programme: full coverage on the unstated grid, extended to the framings by sample.

C.2 Agreement

On the unstated grid: $r = 0.854$, mean bias -0.022 (Opus marginally sterner), 92.1% of paired verdicts within ± 0.5 , 62.8% exact — and the identical five-model ranking, per-tradition tier structure included (fig. 8, Table 9). On the framings sample: $r = 0.777$, bias -0.034 , 94.1% within ± 0.5 , 82.1% exact. Disagreement is graded, not directional: off-diagonal mass sits almost entirely in adjacent half-steps, and sign flips are rare (1.8% of sampled framing cells are a Gemini +1 that Opus scores -1). The one reportable shift is at the top of the unstated ranking: the Sonnet–Inkling tie that is real to Gemini breaks modestly toward Sonnet under Opus (+0.499 vs +0.446 raw-mean over matched cells, where Gemini scored the same cells +0.502 vs +0.493).

On the matched framings sample (Table 10), Opus deflates easy- and medium-tier *guided* scores by ≈ 0.05 – 0.06 and moves the hard tier by -0.01 (guided) and $+0.01$ (stated): the hard-tier guided residual is confirmed almost exactly, and under Opus the guided hard tier sits at $+0.61$ against $+0.87$ medium on the sampled cells — if anything a wider gap than under the primary judge.

C.3 Disclosures

Gemini judged itself in the grid. Its 9,342 own-sittings were not excluded. The Opus layer bounds the effect: on matched unstated cells, Opus scores Gemini’s counsel 0.04 lower than Gemini scores itself (+0.095 vs +0.133) — a real but small self-favor that does not move its rank.

The framings grid was judged safety-on via a router. The unstated grid was judged over the direct API; the stated and guided grids via OpenRouter with safety settings on. A 1,800-judgment pilot put router-vs-direct agreement at $r = 0.930$ (bias $+0.016$) — negligible, but a route change




Figure 8: Dual-judge agreement, cell by cell: joint distribution of paired verdicts (% of cells), Gemini 3.6 Flash × Claude Opus 4.8. Left: full unstated grid. Right: stated+guided sample. Outlined diagonal = exact agreement (62.8% and 82.1%). The sample’s heavier diagonal reflects the stated/guided regime, where verdicts cluster at +1 for both judges.

mid-programme is worth the ink.

The Opus framings sample was collected over two routes. The batch API collected 79% of the sample before a billing cap; the tail was completed live over OpenRouter. 2,597 cells were judged under both routes, giving a same-judge route bridge: $r = 0.949$, bias -0.003 , 98.0% of paired verdicts within ± 0.5 — route equivalence measured, not assumed. Where a cell was judged twice, the later verdict is the one counted.

Table 9: Unstated ranking under each judge: raw pooled means over the 31,114 matched cells (both scopes). This is the validation layer’s scale, not the headline mean-of-tradition-means; rankings, not magnitudes, are the comparable object. The order is identical.

Subject	Gemini 3.6 Flash	Opus 4.8
Claude Sonnet 5	+0.502	+0.499
Inkling	+0.493	+0.446
GPT-5.6 Terra	+0.368	+0.354
Gemini 3.6 Flash	+0.133	+0.095
Qwen3-235B	−0.263	−0.271

Table 10: Tier $\times$ framing on the matched stated+guided sample (post-pressure cells; raw pooled means). Opus deflates near-ceiling guided scores in the easy and medium tiers and confirms the hard tier almost exactly.

Tier	Framing	Opus 4.8	Gemini 3.6	Δ (O−G)	n cells
Easy	Stated	+0.66	+0.66	+0.00	660
Easy	Guided	+0.84	+0.86	−0.03	660
Medium	Stated	+0.59	+0.63	−0.04	630
Medium	Guided	+0.87	+0.94	−0.07	630
Hard	Stated	+0.31	+0.29	+0.02	960
Hard	Guided	+0.58	+0.60	−0.02	960

D Cost and compute

The programme is $519 \times 3 \times 6 \times 5 = 46,710$ sittings. The primary grid is 93,420 judgments (both scopes); the Opus validation layers add 42,711 (31,114 unstated + 9,000 sampled + a 2,597-cell route bridge, batch + live API); the router pilot adds 1,800 — **137,931 judgments** in all. Judges were called concurrently with per-job retry-and-skip; the Gemini grid used prompt caching over the shared rubric and per-scenario guidance, and the Opus layers ran as batch jobs against dedicated capacity. The programme spend:

Table 11: Programme cost (as run). Collection and Gemini-judging figures are priced from the run artifacts at verified 2026-08-03 rates (the framings portion ran in part on funded credits — see Acknowledgements — and its cache accounting makes those figures estimates, not invoices); the unstated Gemini judging is as invoiced on the direct API. The router pilot (1,800 judgments) is included in the Gemini judging figure.

Component	Cost
Subject response collection — five subjects, 46,710 sittings	$\approx$ \$1,106
Judging, Gemini 3.6 Flash — full grid (93,420 judgments)	$\approx$ \$1,317
of which the unstated grid, as invoiced	\$852
Judging, Claude Opus 4.8 — validation layers incl. route bridge (42,711)	$\approx$ \$1,310
Total (as run)	$\approx$ \$3,700

E Additional figures




Figure 9: Every model, unstated $\rightarrow$ stated $\rightarrow$ guided, by tier (per-subject means of tradition means). The dumbbells lengthen left to right across tiers: in the hard tier every model travels most of the scale — Gemini 3.6 Flash from -0.15 to $+0.86$, Qwen3-235B from -0.62 to $+0.19$. Guided endpoints separate the field: three models converge near ceiling in every tier; Terra and Qwen leave visible daylight, in tier order.




Figure 10: Post-pressure score by tradition and subject: filled marker = unstated (with 95% CI), open marker = guided; traditions ordered easiest → hardest unstated. The open markers bunch against the right wall in every tradition except the hard pair — and Qwen’s, which stall mid-scale everywhere. The horizontal distance from filled to open marker is, per cell, the price of not knowing whom you serve.

References

- Ali Abdelaal, Mohammed Nader Al Haffar, Mahmoud Fawzi, and Walid Magdy. IslamicMMLU: A benchmark for evaluating LLMs on islamic knowledge, 2026. arXiv:2603.23750, <https://arxiv.org/abs/2603.23750>.
- Yahyā ibn Sharaf al Nawāwī. *Riyāḍ al-Ṣāliḥīn*. 1270. Compiled c. 1270 CE; consensus-grade hadith compilation read across schools; numerous editions.
- attributed to the Buddha. *Dhammapada*. c. 300 BCE. Khuddaka Nikāya; read across recensions; numerous editions and translations.
- Catholic Church. *Catechism of the Catholic Church*. 1997. Promulgated by Pope John Paul II (*Fidei Depositum*, 1992); editio typica 1997.
- Consortium for Evaluating Faith and Ethics in AI (CEFE-AI). The AllFaith benchmark. <https://cefe.ai>, 2026. Multi-faith evaluation framework (Brigham Young, Baylor, Notre Dame, Yeshiva); datasets at <https://github.com/CEFEAI>.
- Desert Fathers and Mothers. *Apophthegmata Patrum (Systematic Collection)*. 500. Sayings of the Desert Fathers, compiled c. 5th–6th century CE; numerous editions and translations.
- Ezieddin Elmahjub, Junaid Qadir, Abdullah Mushtaq, Rafay Naeem, Ibrahim Ghaznavi, and Waleed Iqbal. IslamicLegalBench: Evaluating LLMs knowledge and reasoning of islamic law across 1,200 years of islamic pluralist legal traditions, 2026. arXiv:2602.21226, <https://arxiv.org/abs/2602.21226>.
- Brett Israelsen, Sheryl Carty, Josh Coates, Nancy Fulda, Julie Park, and Pete Whiting. When AI takes sides on questions of faith: Persistent asymmetries in AI-mediated faith guidance, 2026. arXiv:2605.22975, <https://arxiv.org/abs/2605.22975>.
- M. Waleed Kadous and Benjamin Olsen. JaleesBench: Are AI assistants good spiritual company?, 2026. Companion paper. Code and results: <https://github.com/iaseer-ai/jaleesbench>.
- Abderraouf Lahmar, Md Easin Arafat, Zakarya Farou, and Mufti Mahmud. IslamTrust: A benchmark for LLMs alignment with islamic values. In *5th Muslims in ML Workshop (MusIML), NeurIPS 2025*, 2025. OpenReview: <https://openreview.net/forum?id=PBcv90iKFB>.
- Laozi. *Tao Te Ching (Dào Dé Jīng)*. c. 400 BCE. Traditional attribution; read through the Wang Bi and Heshang Gong commentaries; numerous editions.
- Moshe Chaim Luzzatto. *Mesillat Yesarim*. 1738. The Path of the Upright; classical mussar text; numerous editions.
- MultiWeights. MultiWeights: A model without omissive bias for faiths, 2026. Companion paper (draft; author list forthcoming). Same repository as this work.
- David Wingate, Sheryl Carty, Joshua Coates, Daniel Feldman, Nancy Fulda, Larry Howell, Brett Israelson, Dallin Jacobs, Jonathan Karr, John Paul Kimes, Elisabeth Kincaid, Paul Martens, Gavin Mobley, Suzana Pinheiro, Lindsay Slemboski, and Peter Whiting. Omissive bias in religious representation: Benchmarking LLM answers to everyday ethical decision-making, 2026. AllFaith Religious Representation Benchmark; Consortium for Evaluating Faith and Ethics in AI (CEFE-AI). arXiv:2605.24319, <https://arxiv.org/abs/2605.24319>.