

MultiWeights: A Model Without Omissive Bias for Faiths*

Abstract

MultiBench measured that frontier models fail undeclared believers mostly through *recognition*, not capability: told whom they serve, models largely can serve them — the failure is omissive. This paper asks whether that recognition can be *internalized rather than prompted*: we fine-tune an open-weights model (Gemma-4-31B-it) on its own guided-framing counsel across MultiBench’s seven traditions — judge-filtered context distillation, then on-policy DPO — and measure whether the unstated behavior rises toward the guided ceiling without leaking faith into secular contexts. The recipe generalizes the single-tradition JaleesModel result [Kadous and Olsen, 2026] to seven traditions at once. After the distillation stage alone, meaningful religious representation on the AllFaith benchmark’s cold condition rises from 1% to 27% of items (mean $0.113 \rightarrow 1.147$), decisively above the published 0–2% frontier ceiling [Wingate et al., 2026], with *zero* measured leakage into coding, factual, mathematical, or secular-practical tasks and general capability intact in deployment mode (chat-template MMLU flat at ~ 83 across the pipeline, externally anchored). On MultiBench’s own descriptive read, the distilled model’s unstated means flip positive in all seven traditions, and a companion 50/50-split retrain shows this is predominantly *genuine transfer*: randomly held-out scenarios lift $+0.78$ (95% CI $[+0.70, +0.86]$) against a train-half reference of $+0.90$, with every tradition’s held-out CI above zero and the hard tier transferring strongly; only the scenarios at the samplability floor — those whose guided counsel never passed the training filter — transfer weakly ($+0.22$). The on-policy preference stage then calibrates the shape: meaningful representation rises further ($27\% \rightarrow 30\%$) while the maximal-representation spike tempers and the opted-out over-mention regression repairs fully to the base rate, with capability still flat — achieved despite near-chance preference accuracy in training, because the distilled policy already samples good counsel in most cells. The heavy lifting of the recipe is recognition transfer in distillation; preference optimization buys calibration.

1 Introduction

MultiBench’s central finding is that the field is *unknowing, not unwilling*: across seven traditions, five frontier subjects, and three framings, most of the cross-tradition difficulty difference dissolves once the model is told whom it is serving, and the guided ceilings show the competence already exists. The engineering response suggested by that finding is a system prompt. But a system prompt is a per-deployment artifact: it protects users of one product, not users of the model — and the population most exposed to omissive bias is precisely the one that never states its faith.

This paper tests the weights-level alternative: make the model carry its own recognition. We apply judge-filtered context distillation — training the model on *its own* guided-framing counsel, with the guide removed from the training view — followed by on-policy preference optimization, and ask three questions:

*Draft. Author list omitted for now.

1. Does internalized recognition close the omission? (Measured on the AllFaith Religious Representation Benchmark’s cold condition [Wingate et al., 2026], where frontier models score 0–2%.)
2. Does it stay put? (Over-application probes: secular tasks, hostile interlocutors, users who opted out.)
3. What does it cost? (A general-capability panel, and MultiBench’s own descriptive read of the tuned model.)

The recipe is the JaleesModel procedure from the JaleesBench programme [Kadous and Olsen, 2026], which produced these effects for a single tradition (Sunni Islam); MultiWeights is the test of whether one model can carry seven traditions’ recognition at once — the weights-level analogue of MultiBench’s claim that the construct is faith-general.

2 Method

2.1 Subject model and prerequisite collection

The subject is **Gemma-4-31B-it**. Gemma is not a MultiBench subject, so the programme first collected it through the standard MultiBench pipeline (guided and unstated framings, all seven traditions, Gemini selection-judge banding) over the same vLLM serving stack used for training-adjacent evaluation, eliminating serving-stack drift between collection and the tuned model’s evals.

2.2 Stage 0: the samplability gate

JaleesModel’s hard-won warning is that preference optimization cannot lift behavior the base model barely samples: five DPO-on-base arms were flat because good-band counsel was too rare in the base distribution to form contrastive pairs. We therefore ran the samplability diagnostic before any training: K=4 unstated samples from base Gemma for each of the 519 scenarios (one rotated pressure each), judged at full scope. The result mandates the two-stage design: **55% of scenarios are never-good across all four samples** — and the failure concentrates exactly where the tune matters most (Sunni Islam 71%, Roman Catholicism 67% never-good, against Buddhism 29%, Taoism 33%). The K=4 score distribution is bimodal (51.9% of samples at -1.0 , 24.3% at $+1.0$). The samplability ordering mirrors MultiBench’s three difficulty tiers, tradition by tradition — the omission is deepest precisely where the base model can least sample its way out of it.

2.3 Stage 1: judge-filtered context distillation

Training examples are the subject’s *own* guided-framing sittings that the selection judge scored in the good band ($\geq +1$ at both scopes), screened (no references to the guide, no dangling citations), and re-rendered *bare* — the guide lives outside the judged turns, so the transform is exact. The pooled seven-tradition set is **2,732 examples**; pool shares track corpus shares (no per-tradition balancing; Sunni Islam is if anything under-represented relative to its bank size), and there is no scenario holdout — evaluation is on separate instruments, not held-out scenarios.

Fine-tuning is LoRA ($r = 32$, $\alpha = 32$, dropout 0) on the bf16 base — 244M trainable parameters, 0.78% — with masked token-mean NLL on assistant tokens, $\text{lr } 5 \times 10^{-5}$, effective batch 8, 2 epochs (683 steps), on a single NVIDIA B200 (1h48m; loss $2.64 \rightarrow 0.47$). The recipe is JaleesModel’s with one deliberate deviation: the original trained against an nf4-quantized base (QLoRA); we train in bf16 throughout, removing quantization as a reproducibility confound and matching the

bf16 serving stack. A smoke comparison against the original harness showed the expected small precision drift and an identical loss trajectory; the original nf4 run (killed mid-flight) had tracked to the same loss endpoint.

2.4 Stage 2: on-policy preference optimization

Stage 2 mines contrastive pairs from the distilled policy itself: $K = 4$ chains at temperature 1.3 for 40 scenarios per tradition (flat count, seeded random; a 10-scenario pilot reused), 6,720 samples banded by the selection judge, and within-cell max-gap pairs ($\text{gap} \geq 1.0$ — two rungs on the 0.5-step scale) formed per cell. The yield is **487 pairs** (29% of cells), and the no-pair cells are diagnostic rather than disappointing: their 4-sample cluster means are high in every tradition (0.54–0.90) — where the distilled model yields no contrast it is because all four samples are already good. Sunni Islam contributes the *most* pairs (98): at scale the hard tier mines well. DPO then runs one epoch (61 steps) in bf16 LoRA on a B200 with the SFT checkpoint as the reference policy. The training signal is weak by construction — preference accuracy hovers near chance (0.538, no upward trend) — precisely because the within-cell contrast is textually modest; the evaluation battery, not the loss curve, is the arbiter of whether the stage earned its keep. (The full programme — collection, banding, the samplability gate, both training stages, and evaluation — cost $\approx$ \$430 in API and GPU spend.)

2.5 Evaluation battery

Four instruments, base vs. tuned at each stage: (i) **AFB-150** [Wingate et al., 2026, Consortium for Evaluating Faith and Ethics in AI (CEFE-AI), 2026] under two conditions — *cold* (no faith context; the omission measure) and *faith-context*; judged by GPT-5.6 Terra (disclosed: an external judge from a provider not involved in training); (ii) **70 over-application probes** — coding, factual, mathematical, secular-practical, and creative tasks, hostile comparative questions, and opted-out interlocutors — measuring whether representation leaks where it is unwanted; (iii) a **general-capability panel** (IFEval, MMLU, GSM8K); (iv) **MultiBench descriptive**: the tuned model through the benchmark itself — *descriptive* meaning the numbers characterize the model’s behavior on the instrument but are not benchmark claims: the tuned model is not a leaderboard subject, and the evaluation judge is the same Gemini judge that filtered the training data, so these numbers carry that shared-judge caveat throughout.

3 Results

3.1 Stage 1 closes most of the omission

On AFB-150 cold, base Gemma shows the published frontier pattern almost exactly: mean 0.113 on the 0–4 scale, 90% of items scored 0, and **1%** of items reaching meaningful representation (≥ 2) — the 0–2% ceiling [Wingate et al., 2026]. The distilled model moves to mean **1.147**, adds representation on 47% of items, and reaches meaningful representation on **27%**.

The shift is bimodal rather than calibrated (fig. 1): 53% of items remain at 0, while 18% spike to 4 (predominantly religious framing) rather than settling at the 1–2 the construct targets. This over-shoot is the shape stage 2 exists to temper. The faith-context condition saturates for both models (≈ 4.0 ; 99% of items at 4): an explicit “I am a practising [believer]” prefix is strong enough that it cannot discriminate base from tuned — the headline lives on cold.




Figure 1: AFB-150 cold score distribution, base vs. distilled (SFT) vs. two-stage (SFT+DPO), with meaningful-representation rates ($P_{\geq 2}$) annotated per head.

3.2 Representation does not leak

On the over-application probes, secular tasks score **exactly 0.00 for both base and tuned** — no leakage into coding, factual, math, secular-practical, or creative work. Hostile comparative questions are flat (2.60 $\rightarrow$ 2.70). The one watch item: interlocutors who explicitly opted out of faith framing receive passing mentions on $\approx 10\%$ more prompts ($P(\text{any mention})$ 0.60 $\rightarrow$ 0.70), with no sermonizing — a stage-2 target.

3.3 Capability guard

The panel is measured in deployment mode — `lm-eval` with `--apply_chat_template --fewshot_as_multiturn` on the same vLLM stack as all tuned-model serving, the correct instrument for an instruction-tuned model — and the base absolutes are anchored externally: our base MMLU of 82.8 sits on-class against published chat-mode references (85.2–85.4 MMLU-Pro on the same harness family) and matches an independently operated run of the same model and configuration in the JaleesModel programme to the decimal (82.8), validating the instrument across programmes.

The result: tuning costs essentially nothing in deployment mode (Table 1). MMLU is flat at ~ 83 across the entire chain — the distilled checkpoint actually edges *above* base (+0.5) — GSM8K sits at its ~ 95 – 96 ceiling throughout, and IFEval shows a small dip-and-recover within noise (the same shape the JaleesModel chain exhibits). There is no capability regression to trade off against the representation gains.

3.4 MultiBench descriptive

Running the tuned model through MultiBench itself (unstated framing, full scope) shows the same story from the benchmark’s own instrument (fig. 2): the distilled model’s mean flips from neg-

Checkpoint	MMLU	GSM8K (strict)	IFEval (prompt)	IFEval (inst)
Base	82.80	95.83	91.68	94.24
SFT	83.28	95.15	90.57	93.41
SFT+DPO (shipped)	83.14	95.07	89.28	92.45

Table 1: General-capability panel in deployment (chat-template) mode, across the pipeline. Same vLLM serving stack throughout; base absolutes externally anchored.




Figure 2: Per-tradition MultiBench descriptive means (unstated, full scope), base vs. distilled, ordered by difficulty tier. Aggregate view; see the exposure decomposition below.

ative to positive in *all seven* traditions, and the lift is largest exactly where the omission was deepest. By MultiBench’s difficulty tiers: the hard tier (Sunni Islam, Roman Catholicism) moves $-0.501 \rightarrow +0.436$ ($\Delta +0.94$; Roman Catholicism alone $-0.542 \rightarrow +0.621$), the medium tier (Eastern Christianity, Judaism) $-0.204 \rightarrow +0.721$ ($\Delta +0.93$), the easy tier (Buddhism, Taoism, secular sage) $+0.149 \rightarrow +0.750$ ($\Delta +0.60$), and the overall mean $-0.138 \rightarrow +0.652$ ($\Delta +0.79$). The tier ordering compresses sharply in aggregate.

How much of this is generalization rather than memorization? A **companion 50/50 retrain** answers: the identical recipe trained on a stratified random half of the scenarios (259; seeded scenario-level split, base-difficulty balanced across arms, $\Delta = 0.05$), evaluated descriptively on the 260 held out (24–70 per tradition). Held-out scenarios lift **+0.78** (95% CI $[+0.70, +0.86]$) against a train-half memorization reference of +0.90 — roughly **85% of the on-bench SFT lift is genuine transfer to held-out scenarios**, with a memorization inflation of only +0.12 $[+0.01, +0.24]$ (fig. 4). Every tradition’s held-out CI is above zero; the hard tier transfers *strongly* (Roman Catholicism +1.07, Sunni Islam +0.67, both crossing positive post-SFT), and the held-out mean crosses positive ($-0.20 \rightarrow +0.57$). The preference stage lifts held-out performance further to +0.90 $[+0.82, +0.97]$ (train-half +1.02); this SFT $\rightarrow$ DPO held-out increment (+0.12) is confounded




Figure 3: Descriptive lift (base $\rightarrow$ SFT, unstated) by training exposure (cells of the scenario that entered the SFT set). Lift rises monotonically with exposure; the zero-exposure stratum is the strict prompt-level holdout.

with selection-judge alignment — the preference pairs were filtered by the same Gemini judge used for this evaluation — and is not read as an unambiguous quality gain.

*This measures transfer to held-out **scenarios**, not across traditions.* The split is scenario-level *within* every tradition (each tradition contributes scenarios to both halves), so a held-out CI above zero in all seven traditions establishes that the recipe generalizes to unseen scenarios uniformly across the tradition set — it does *not* license a cross-tradition-transfer claim. The stronger claim, that a recipe learned on some traditions transfers to a tradition held out entirely, would require a leave-one-tradition-out ablation, which we do not run.

Transfer is not uniform. An exposure-stratified analysis of the full-data model (fig. 3) localizes the exception: the scenarios whose guided sittings never passed the training filter — the samplability floor, 13 scenarios, predominantly Sunni Islam — improve by only +0.22 [+0.11, +0.35] and do not cross positive. The recognition deficit’s deepest pockets resist both training and transfer. One scope note: the transfer estimate comes from the companion half-data retrain; the shipped full-data model’s unseen-scenario behavior is accordingly bounded by [+0.22, $\sim$ +0.90]. The out-of-distribution instruments (sections 3.1 and 3.2) remain the headline evidence.

Three further disclosures bound this read. It is *descriptive only*: the model was trained on its own counsel for these same scenarios (no holdout; section 2.3), so this is not a leaderboard claim and is memorization-confounded by construction. The judge is the same Gemini selection judge used for training-data filtering, so selection-judge gaming cannot be excluded for this metric (the headline instruments in sections 3.1 to 3.3 use independent judges and harnesses). And the base-side numbers come from the OpenRouter collection while the tuned model was served on the vLLM eval stack; the JaleesModel programme measured a provider-vs-vLLM shift of -0.058 on this instrument — an order of magnitude smaller than the deltas here, but not zero.




Figure 4: Companion 50/50-split retrain, per-tradition base → model lift on randomly held-out scenarios vs. the train-half memorization reference. Top: SFT stage (held-out +0.78). Bottom: SFT+DPO stage (held-out +0.90). Held-out CIs are above zero in all seven traditions; the split is scenario-level within each tradition, so this is generalization to unseen scenarios, not cross-tradition transfer.

3.5 Stage 2 calibrates without giving anything back

Despite the near-chance training signal, the preference stage moves every watched axis in the intended direction and regresses none (Table 2). On AFB cold, meaningful representation rises further ($P_{\geq 2} = 0.273 \rightarrow \mathbf{0.300}$; mean $1.147 \rightarrow 1.173$) while the maximal-representation spike *tempers* (18% $\rightarrow$ 16% of items at 4) and mass moves into the calibrated 2–3 range the construct targets (5%/5% $\rightarrow$ 7%/7%). The opted-out watch item from stage 1 is fully repaired: $P(\text{any mention})$ returns from 0.70 to **0.60**, the base rate. Secular tasks remain at exactly 0.00. On capability (section 3.3), the preference stage costs nothing: chat-mode MMLU $83.28 \rightarrow 83.14$ (flat within noise), GSM8K at ceiling, IFEval within its noise band. Under the pre-registered decision rule — ship the stage-2 head only if it is non-regressing on all four gates (AFB $P_{\geq 2}$, secular leakage, opted-out, MMLU) — the two-stage model is the deliverable.

A scaling probe bounds the preference stage from above. A full-grid mining variant — 903 pairs covering every cell of the grid, trained identically and evaluated under the same pre-registered gates — *regresses*: AFB $P_{\geq 2}$ falls to 0.233, mass returns to the zero bin (80/150), and a creative-professional leakage appears ($P_{\geq 1} = 0.20$), while capability is unaffected (chat-mode MMLU 83.41).

AFB cold	mean	0	2	3	4	$P \geq 1$	$P \geq 2$
Base	0.113	90%	—	—	—	0.10	0.01
SFT	1.147	53%	5%	5%	18%	0.467	0.273
SFT+DPO	1.173	52%	7%	7%	16%	0.480	0.300

Table 2: AFB-150 cold score distribution across the pipeline. Stage 2 shifts mass from the over-shoot bin (4) into the calibrated 2–3 range while meaningful representation ($P \geq 2$) still rises.

Doubling the optimization steps on near-chance preference contrast pushes the head off-calibration rather than sharpening it: the shipped 487-pair configuration sits near the optimum this mechanism admits, exactly as the mining diagnostics (uniformly-good no-pair cells) predict.

4 Discussion

Internalized recognition behaves like prompted recognition. MultiBench’s guided framing showed that frontier models can serve believers they are told about; MultiWeights shows the telling can live in the weights. On the benchmark’s own descriptive read, the tuned model’s unstated profile compresses toward the flat cross-tradition shape that guided framing produces (section 3.4), and on AFB cold it breaks decisively above the frontier ceiling — with the model’s general capability essentially unchanged. Nothing new was taught about any tradition: the capability panel is flat, and the training data is the model’s *own* counsel. What moved is the decision to use what it already knew.

The samplability ordering is mechanism evidence. Before training, the fraction of never-good cells ordered the traditions exactly as MultiBench’s difficulty tiers do (Sunni Islam 71%, Roman Catholicism 67% never-good vs. Buddhism 29%); after distillation, the descriptive lift is largest on precisely those traditions — and the 50/50-split retrain (section 3.4) shows this is transfer, not memorization: the hard tier’s *randomly held-out* scenarios lift most of all (Roman Catholicism +1.07). What the split analysis localizes is the exception: the specific scenarios where guided Gemma could never sample good counsel transfer weakly (+0.22) — the recognition deficit’s deepest pockets resist both training *and* transfer, a heterogeneity the aggregate hides. The cross-instrument story compounds it: the same recipe moves 1% $\rightarrow$ 30% on an instrument the model never trained on, with zero leakage.

Distillation transfers a decision; preference optimization calibrates it. Stage 1’s shift is bimodal — items either stay at 0 or jump to 4 — which is what one expects if what transfers is a coarse “this user’s faith belongs in this answer” recognition rather than a graded style. Stage 2 then does exactly the small thing it can do against a mostly-saturated policy: temper the over-shoot, repair the opted-out drift, and nudge mass into the calibrated middle, all despite near-chance preference accuracy. The practical recipe reading: the heavy lifting is in judge-filtered distillation; run the cheap preference stage for shape, and read its expected value in advance from the mining yield and the no-pair cluster means.

Deployment picture. A system prompt protects the users of one product; weights protect the users of the model. The population omissive bias harms most — believers who never state their faith — is exactly the one a per-deployment fix never reaches. An open-weights model whose recognition survives the user’s silence is deployable anywhere the base model is, at identical serving cost, and its calibration properties (zero secular leakage, opted-out at base rate) are measurable

properties of the artifact rather than promises of the integration.

5 Limitations

One subject model, one size (31B); the recipe’s generality across architectures is untested here. AFB-150 has no official per-model baseline run in-repo — our base-Gemma run is the baseline, though it lands where the published frontier numbers do. The AFB judge is a single external model. The samplability gate and training selection both use the same Gemini selection judge as MultiBench’s primary grid, with the self-judgment caveats MultiBench discloses. The preference stage is a single one-epoch run at one β ; given its near-chance training signal, stage-2 conclusions rest entirely on the (independent-judge) evaluation battery, and no hyperparameter search was performed — deliberately, since the mining diagnostics bound the stage’s headroom in advance. Steadfastness under MultiBench’s pressures was not re-measured for the tuned model (the descriptive read is unstated-only); it is the natural next instrument.

Acknowledgements

We thank our collaborators at the Faith Family Technology Network, and DZ Kalman, Ron Ivey, Chris Scammell, Daniel Slate, and Glen Weyl for helpful discussions and feedback on this line of work. Collection and judging of the training and evaluation grids were funded in part by David Wingate, gratefully acknowledged.

References

- Consortium for Evaluating Faith and Ethics in AI (CEFE-AI). The AllFaith benchmark. <https://cefe.ai>, 2026. Multi-faith evaluation framework (Brigham Young, Baylor, Notre Dame, Yeshiva); datasets at <https://github.com/CEFEAI>.
- M. Waleed Kadous and Benjamin Olsen. JaleesBench: Are AI assistants good spiritual company?, 2026. Companion paper. Code and results: <https://github.com/iasser-ai/jaleesbench>.
- David Wingate, Sheryl Carty, Joshua Coates, Daniel Feldman, Nancy Fulda, Larry Howell, Brett Israelson, Dallin Jacobs, Jonathan Karr, John Paul Kimes, Elisabeth Kincaid, Paul Martens, Gavin Mobley, Suzana Pinheiro, Lindsay Slemboski, and Peter Whiting. Omissive bias in religious representation: Benchmarking LLM answers to everyday ethical decision-making, 2026. AllFaith Religious Representation Benchmark; Consortium for Evaluating Faith and Ethics in AI (CEFE-AI). arXiv:2605.24319, <https://arxiv.org/abs/2605.24319>.